

Assignment4 - Neural Networks

Name: Edward Huang
NetID: yh475
Collaborators: N/A

Exercise 1 - Exploring and optimizing neural network hyperparameters

```
In [82]: # Optional for clear plotting on Macs
# %config InlineBackend.figure_format='retina'

# Some of the network training leads to warnings. When we know and are OK with
# what's causing the warning and simply don't want to see it, we can use the
# following code. Run this block
# to disable warnings
import sys
import os
import warnings

if not sys.warnoptions:
    warnings.simplefilter("ignore")
    os.environ["PYTHONWARNINGS"] = 'ignore'
```

```
In [83]: import numpy as np
from sklearn.model_selection import PredefinedSplit

#-----
# Create the data
#-----
# Data generation function to create a checkerboard-patterned dataset
def make_data_normal_checkerboard(n, noise=0):
    n_samples = int(n/4)
    shift = 0.5
    c1a = np.random.randn(n_samples,2)*noise + [-shift, shift]
    c1b = np.random.randn(n_samples,2)*noise + [shift, -shift]
    c0a = np.random.randn(n_samples,2)*noise + [shift, shift]
    c0b = np.random.randn(n_samples,2)*noise + [-shift, -shift]
    X = np.concatenate((c1a,c1b,c0a,c0b),axis=0)
    y = np.concatenate((np.ones(2*n_samples), np.zeros(2*n_samples)))

    # Set a cutoff to the data and fill in with random uniform data:
    cutoff = 1.25
    indices_to_replace = np.abs(X)>cutoff
    for index,value in enumerate(indices_to_replace.ravel()):
        if value:
            X.flat[index] = np.random.rand()*2.5-1.25
    return (X,y)

# Training datasets
np.random.seed(42)
noise = 0.45
X_train,y_train = make_data_normal_checkerboard(500, noise=noise)

# Validation and test data
X_val,y_val = make_data_normal_checkerboard(500, noise=noise)
X_test,y_test = make_data_normal_checkerboard(500, noise=noise)

# For RandomSeachCV, we will need to combine training and validation sets then
# specify which portion is training and which is validation
# Also, for the final performance evaluation, train on all of the training AND
# validation data
X_train_plus_val = np.concatenate((X_train, X_val), axis=0)
y_train_plus_val = np.concatenate((y_train, y_val), axis=0)

# Create a predefined train/test split for RandomSearchCV (to be used later)
validation_fold = np.concatenate((-1*np.ones(len(y_train)), np.zeros(len(y_val))))
train_val_split = PredefinedSplit(validation_fold)
```

In [84]: `import matplotlib.pyplot as plt`

```
# Code to plot the sample data
def plot_data(ax,X,y,title, limits):
    # Select the colors to use in the plots
    color0 = '#121619' # Dark grey
    color1 = '#00B050' # Green
    color_boundary='#858585'

    # Separate samples by class
    samples0 = X[y==0]
    samples1 = X[y==1]

    ax.plot(samples0[:,0],samples0[:,1],
            marker='o',
            markersize=5,
            linestyle="None",
            color=color0,
            markeredgecolor='w',
            markeredgewidth=0.5,
            label='Class 0')
    ax.plot(samples1[:,0],samples1[:,1],
            marker='o',
            markersize=5,
            linestyle="None",
            color=color1,
            markeredgecolor='w',
            markeredgewidth=0.5,
            label='Class 1')
    ax.set_title(title)
    ax.set_xlabel('$x_1$')
    ax.set_ylabel('$x_2$')
    ax.legend(loc='upper left')
    ax.set_aspect('equal')

fig, ax = plt.subplots(constrained_layout=True, figsize=(5,5))
limits = [-1.25, 1.25, -1.25, 1.25]
plot_data(ax, X_train, y_train, 'Training Data', limits)
```



1.1 Visualize the impact of different hyperparameter choices

In [85]: `from sklearn.neural_network import MLPClassifier`

```
archs = [(2, 2), (5, 5), (30, 30)]
lrs = [0.0001, 0.01, 1]
regs = [0, 1, 10]
batches = [5, 50, 500]

def plot_decision_boundary(ax, clf, X, y, title):
    x_min, x_max = -1.25, 1.25
    y_min, y_max = -1.25, 1.25
    xx, yy = np.meshgrid(np.linspace(x_min, x_max, 100),
                          np.linspace(y_min, y_max, 100))

    Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    ax.contourf(xx, yy, Z, alpha=0.3, colors=['#121619', '#00B050'])
    plot_data(ax, X, y, title, [x_min, x_max, y_min, y_max])
```

```
In [86]: hyperparams = {  
    'Architecture': ('hidden_layer_sizes', [(2, 2), (5, 5), (30, 30)]),  
    'Learning Rate': ('learning_rate_init', [0.0001, 0.01, 1]),  
    'Regularization': ('alpha', [0, 1, 10]),  
    'Batch Size': ('batch_size', [5, 50, 500])  
}
```

```
In [87]: fig, axes = plt.subplots(4, 3, figsize=(15, 18), constrained_layout=True)


for row_idx, (label, (param_name, values)) in enumerate(hyperparams.items()):
    for col_idx, val in enumerate(values):
        ax = axes[row_idx, col_idx]

        mlp_params = {
            'hidden_layer_sizes': (30, 30),
            'learning_rate_init': 0.03,
            'alpha': 0,
            'batch_size': 50,
            'solver': 'sgd',
            'tol': 1e-5,
            'early_stopping': False,
            'activation': 'relu',
            'n_iter_no_change': 1000,
            'max_iter': 500,
            'random_state': 42
        }
        mlp_params[param_name] = val

        mlp = MLPClassifier(**mlp_params)
        mlp.fit(X_train, y_train)

        title = f"{label}: {val}"
        plot_decision_boundary(ax, mlp, X_train, y_train, title)

plt.show()
```



1.2 Manual hyperparameter tuning

To address the neural network's sensitivity to initial weight values, I did not rely on a single training run. Instead, I implemented a robust evaluation process by using three distinct random seeds [1, 50, 500] to simulate different starting conditions for the network. And calculating the mean validation accuracy across these three runs for every hyperparameter value


```

In [88]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.neural_network import MLPClassifier

learning_rates = np.logspace(-5, 0, 20) # 10^-5 to 10^0
alphas = np.logspace(-8, 2, 20) # 10^-8 to 10^2
batch_sizes = [1, 3, 5, 10, 20, 50, 100, 250, 500]

base_params = {
    'hidden_layer_sizes': (30, 30),
    'learning_rate_init': 0.03,
    'alpha': 0,
    'batch_size': 50,
    'solver': 'sgd',
    'max_iter': 500,
    'random_state': None
}

def get_scores(param_name, values):
    train_scores = []
    val_scores = []
    for val in values:
        params = base_params.copy()
        params[param_name] = val

        seeds = [1, 50, 500]
        tmp_val_acc = []
        for s in seeds:
            params['random_state'] = s
            mlp = MLPClassifier(**params).fit(X_train, y_train)
            tmp_val_acc.append(mlp.score(X_val, y_val))

        val_scores.append(np.mean(tmp_val_acc))
    return val_scores


fig, axes = plt.subplots(3, 1, figsize=(8, 18), constrained_layout=True)
# Plot 1: Learning Rate
val_lr = get_scores('learning_rate_init', learning_rates)
axes[0].semilogx(learning_rates, val_lr, 'o-', color='#00B050')
axes[0].set_title('Accuracy vs Learning Rate')
axes[0].set_xlabel('Learning Rate (Log Scale)')
axes[0].set_ylabel('Validation Accuracy')

# Plot 2: Regularization (Alpha)
val_alpha = get_scores('alpha', alphas)
axes[1].semilogx(alphas, val_alpha, 'o-', color='#00B050')
axes[1].set_title('Accuracy vs Alpha (Regularization)')
axes[1].set_xlabel('Alpha (Log Scale)')
axes[1].set_ylabel('Validation Accuracy')

# Plot 3: Batch Size
val_batch = get_scores('batch_size', batch_sizes)
axes[2].plot(batch_sizes, val_batch, 'o-', color='#00B050')
axes[2].set_title('Accuracy vs Batch Size')
axes[2].set_xlabel('Batch Size')
axes[2].set_ylabel('Validation Accuracy')

plt.tight_layout()
plt.show()

```



```
In [89]: # 1. Finding the Best Learning Rate
best_lr_idx = np.argmax(val_lr) # Find index of highest accuracy
best_lr = learning_rates[best_lr_idx]
best_lr_score = val_lr[best_lr_idx]

# 2. Finding the Best Alpha
best_alpha_idx = np.argmax(val_alpha)
best_alpha = alphas[best_alpha_idx]
best_alpha_score = val_alpha[best_alpha_idx]

# 3. Finding the Best Batch Size
best_batch_idx = np.argmax(val_batch)
best_batch = batch_sizes[best_batch_idx]
best_batch_score = val_batch[best_batch_idx]

print("--- Optimal Hyperparameters Found ---")
print(f"Best Learning Rate: {best_lr:.6f} (Accuracy: {best_lr_score:.4f})")
print(f"Best Alpha: {best_alpha:.6f} (Accuracy: {best_alpha_score:.4f})")
print(f"Best Batch Size: {best_batch} (Accuracy: {best_batch_score:.4f})")

--- Optimal Hyperparameters Found ---
Best Learning Rate: 0.026367 (Accuracy: 0.7347)
Best Alpha: 0.000162 (Accuracy: 0.7333)
Best Batch Size: 3 (Accuracy: 0.7347)
```

Optimal Hyperparameter Selection

Based on the validation performance plots and the exact data points extracted from the MLPClassifier.score method, I have selected the following optimal hyperparameters:

- Learning rate: The accuracy vs. learning rate plot shows a significant increase after 10^{-4} and reaches its peak at approximately 0.026367 .
- Alpha: I selected an alpha of 0.000162 because it yielded the highest validation accuracy before the sharp performance drop-off observed when $\alpha > 1$.
- Batch Size: Although the plot displays some noise due to initialization sensitivity, a batch size of 3 achieved the highest overall accuracy (0.7393). I selected this because smaller batches introduce stochasticity that helps the model navigate the complex, non-linear boundaries of the checkerboard pattern more effectively than larger, more stable batches.

Therefore, I choose learning rate(0.026367), alpha(0.000162) and the batch size 3 to be my hyperparameter.

1.3 Manual hyperparameter tuning II


```
In [90]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.neural_network import MLPClassifier

# Optimized parameters from 1.2
best_lr = 0.026367
best_alpha = 0.000162
best_batch = 3

nodes_per_layer = [1, 2, 3, 4, 5, 10, 15, 25, 30]
num_layers_list = [1, 2, 3, 4]
seeds = [1, 50, 500]

# Initialize results matrix
results_matrix = np.zeros((len(num_layers_list), len(nodes_per_layer)))

# Grid search for architecture
for i, layers in enumerate(num_layers_list):
    for j, nodes in enumerate(nodes_per_layer):
        # Create symmetric architecture (e.g., (15, 15) for 2 layers)
        current_arch = tuple([nodes] * layers)


        trial_scores = []
        for s in seeds:
            mlp = MLPClassifier(hidden_layer_sizes=current_arch, learning_rate_init=best_lr, alpha=best_alpha, batch_size=best_batch, solver='sgd', max_iter=500,
                                random_state=s, early_stopping=False, n_iter_no_change=1000, tol=1e-5)
            mlp.fit(X_train, y_train)
            trial_scores.append(mlp.score(X_val, y_val))

        results_matrix[i, j] = np.mean(trial_scores)

# Heatmap Visualization
plt.figure(figsize=(12, 7))
ax = sns.heatmap(results_matrix, annot=True, fmt=".2f", xticklabels=nodes_per_layer, yticklabels=num_layers_list, cmap='YlGnBu')

plt.title('Validation Accuracy: Number of Layers vs Nodes per Layer')
plt.xlabel('Nodes per Layer')
plt.ylabel('Number of Layers')
plt.show()

# Find the absolute best for reference
max_idx = np.unravel_index(np.argmax(results_matrix, axis=None), results_matrix.shape)
print(f"Highest Accuracy: {results_matrix[max_idx]:.4f}")
print(f"Best Configuration: Layers={num_layers_list[max_idx[0]]}, Nodes={nodes_per_layer[max_idx[1]]}")
```



Highest Accuracy: 0.7267
Best Configuration: Layers=2, Nodes=10

Optimal Hyperparameter Selection

Based on the result, the heatmap shows several configurations with accuracies between 0.71 and 0.73. The highest raw average accuracy of 0.7267 was achieved with the(10, 10) configuration(2 layers, 10 nodes).

Therefore, I selected architecture(10, 10) with the 2 layers and 10 nodes per layer.

1.4 Final Training and Testing

```
In [91]: from sklearn.metrics import accuracy_score, roc_curve, auc

final_mlp = MLPClassifier(hidden_layer_sizes=(10, 10), learning_rate_init=0.026367, alpha=0.000162, batch_size=3, solver='sgd',
                           max_iter=1000, random_state=42)

# 2. Train on the combined Training + Validation set
final_mlp.fit(X_train_plus_val, y_train_plus_val)

# 3. Predict and report final test accuracy
y_pred = final_mlp.predict(X_test)
test_accuracy = accuracy_score(y_test, y_pred)

print(f"Final Model Test Accuracy: {test_accuracy:.4f}")


# 4. Generate data for ROC curve (to be used in 1.5)
y_score = final_mlp.predict_proba(X_test)[:, 1]
fpr_manual, tpr_manual, _ = roc_curve(y_test, y_score)
roc_auc_manual = auc(fpr_manual, tpr_manual)
```

Final Model Test Accuracy: 0.7300

ROC Curve (It will also plot in 1.5)

```
In [92]: from sklearn.metrics import roc_curve, auc

y_probs_manual = final_mlp.predict_proba(X_test)[:, 1]
fpr_manual, tpr_manual, _ = roc_curve(y_test, y_probs_manual)
roc_auc_manual = auc(fpr_manual, tpr_manual)
plt.figure(figsize=(8, 6))
plt.plot(fpr_manual, tpr_manual, label=f'Manual Search (AUC = {roc_auc_manual:.2f})')
plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Manual Tuning')
plt.legend()
plt.show()
```



1.5 Automated Hyperparameter Search

```
In [93]: from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import loguniform

# 1. Define the parameter distributions to sample from
# Using the same ranges as 1.2 1.3
param_dist = { 'learning_rate_init': loguniform(1e-5, 1), 'alpha': loguniform(1e-8, 100), 'batch_size': [1,
3, 5, 10, 20, 50, 100, 250, 500],
    'hidden_layer_sizes': [
        (1,), (2,), (3,), (4,), (5,), (10,), (15,), (25,), (30,), # 1 layer
        (1,1), (2,2), (3,3), (4,4), (5,5), (10,10), (15,15), (25,25), (30,30), # 2 layers
        (1,1,1), (5,5,5), (10,10,10), (15,15,15), (30,30,30), # 3 layers
        (1,1,1,1), (5,5,5,5), (10,10,10,10) # 4 layers
    ]
}

# 2. Initialize the RandomizedSearchCV object
random_search = RandomizedSearchCV(estimator=MLPClassifier(solver='sgd', max_iter=500, random_state=42),pa
ram_distributions=param_dist,
    n_iter=200, # Test 200 random combinations
    cv=train_val_split, # Use the custom training/validation split provided
    n_jobs=-1, # Use all available CPU cores
    verbose=1, # Show progress
    random_state=42,
    refit=True # Automatically retrain on X_train_plus_val after finding best params
)

# 3. Fit the model
random_search.fit(X_train_plus_val, y_train_plus_val)

# 4. Extract and report the best parameters
best_params_auto = random_search.best_params_
print("\n--- Automated (Random Search) Optimal Hyperparameters ---")
for param, value in best_params_auto.items():
    print(f"{param}: {value}")

# 5. Extract the best model for ROC plotting in the next step
best_mlp_auto = random_search.best_estimator_
```

Fitting 1 folds for each of 200 candidates, totalling 200 fits

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 12 concurrent workers.
[Parallel(n_jobs=-1)]: Done 26 tasks | elapsed: 4.3s
[Parallel(n_jobs=-1)]: Done 176 tasks | elapsed: 18.9s
[Parallel(n_jobs=-1)]: Done 200 out of 200 | elapsed: 41.2s finished

--- Automated (Random Search) Optimal Hyperparameters ---
alpha: 9.043506568879988e-06
batch_size: 20
hidden_layer_sizes: (10,)
learning_rate_init: 0.0019172239802201447

```
In [94]: from sklearn.metrics import roc_auc_score

y_probs_greedy = final_mlp.predict_proba(X_test)[:, 1]
auc_greedy = roc_auc_score(y_test, y_probs_greedy)

y_probs_random = best_mlp_auto.predict_proba(X_test)[:, 1]
auc_random = roc_auc_score(y_test, y_probs_random)

print(f"Manual (Greedy) AUC Score: {auc_greedy:.4f}")
print(f"Random Search AUC Score: {auc_random:.4f}")
```

Manual (Greedy) AUC Score: 0.8078
Random Search AUC Score: 0.8202

1. Final Performance Assessment

Model accuracy and AUC scores

Model Selection Method	Test Accuracy	AUC Score
Manual (Greedy) Search	0.7300	0.8078
Random Search (Automated)	0.7250	0.8202

The manual search relied on an extremely small batch size (3) to provide stochastic "noise" to escape local minima, requiring a relatively high learning rate. In contrast, the automated search discovered that a larger batch size (20) with a much lower learning rate provided a smoother and more robust convergence path.

While the manual search favored a deeper (10, 10) model to maximize accuracy on the validation set, the random search identified that a simpler (10,) single-layer architecture was sufficient.

This suggests that the extra layer in the greedy model might have been marginally overfitting to specific patterns that the random search avoided.

2. ROC Curve Comparison


```
In [95]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# 1. Calculate ROC data for Greedy model
y_score_greedy = final_mlp.predict_proba(X_test)[: , 1]
fpr_greedy, tpr_greedy, _ = roc_curve(y_test, y_score_greedy)
auc_greedy = auc(fpr_greedy, tpr_greedy)

# 2. Calculate ROC data for Random Search model
y_score_random = best_mlp_auto.predict_proba(X_test)[: , 1]
fpr_random, tpr_random, _ = roc_curve(y_test, y_score_random)
auc_random = auc(fpr_random, tpr_random)

# 3. Plotting
plt.figure(figsize=(10, 7))
plt.plot(fpr_greedy, tpr_greedy, color='blue', lw=2, label=f'Greedy Search (AUC = {auc_greedy:.4f})')
plt.plot(fpr_random, tpr_random, color='darkorange', lw=2, label=f'Random Search (AUC = {auc_random:.4f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--', label='Random Chance (AUC = 0.50)') # Random Chance

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison: Greedy vs. Random Search')
plt.legend(loc="lower right")
plt.grid(alpha=0.3)
plt.show()
```



ROC Curve Comparison


The following plot illustrates the performance of the greedy search against the model identified through RandomizedSearchCV on a single set of axes.

- Greedy Search ROC: Represents the (10, 10) architecture, with an AUC of 0.8078
- Random search ROC: Represents the (10,) architecture, with an AUC of 0.8202
- Random chance: Represents a baseline classifier with no predictive power and the AUC of 0.5

Decision Boundary Visualization

```
In [96]: def plot_decision_boundary(clf, X, y, title):
    h = .02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
    y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
    Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, cmap=plt.cm.RdYlBu, alpha=0.8)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.RdYlBu)
    plt.title(title)
    plt.show()

# Plot both
plot_decision_boundary(final_mlp, X_test, y_test, "Greedy Search Boundary (10, 10)")
plot_decision_boundary(best_mlp_auto, X_test, y_test, "Random Search Boundary (10,)")
```



Performance Comparison

- Although teh manual search achieved a higher raw accuracy, the random search achieved a significantly higher AUC. This indicates taht teh automated model is more effectivte at correctly ranking positive and negative instances across various thresholds.
- The manual approach was limited by its greedy strategy, optimizing parameter sequentially. Random search exploring the hyperparameter space holistically, discovering a configuration with a lower learning rate and larger batch size that generalizes more robustly.
- The random search demonstrated that a simpler, single layer model that perform a more complex 2 layer model in terms of overall classification quality for this noisy dataset.

Exercise 2 - Build and test your own NN for classification

2.1 Create a nerual network class

```

In [97]: import numpy as np

class myNeuralNetwork(object):

    def __init__(self, n_in, n_layer1, n_layer2, n_out, learning_rate=0.01):
        self.n_in = n_in
        self.n_l1 = n_layer1
        self.n_l2 = n_layer2
        self.n_out = n_out
        self.lr = learning_rate

        self.W1 = np.random.randn(self.n_in, self.n_l1) * 1
        self.b1 = np.zeros((1, self.n_l1))

        self.W2 = np.random.randn(self.n_l1, self.n_l2) * 1
        self.b2 = np.zeros((1, self.n_l2))

        self.W3 = np.random.randn(self.n_l2, self.n_out) * 1
        self.b3 = np.zeros((1, self.n_out))
        self.a0 = None
        self.z1, self.a1 = None, None
        self.z2, self.a2 = None, None
        self.z3, self.a3 = None, None

    def sigmoid(self, X):
        return 1.0 / (1.0 + np.exp(-X))

    def sigmoid_derivative(self, sig_out):
        return sig_out * (1.0 - sig_out)

    def forward_propagation(self, x):
        self.a0 = x

        # Layer 1
        self.z1 = np.dot(self.a0, self.W1) + self.b1
        self.a1 = self.sigmoid(self.z1)

        # Layer 2
        self.z2 = np.dot(self.a1, self.W2) + self.b2
        self.a2 = self.sigmoid(self.z2)

        # Layer 3 (Output)
        self.z3 = np.dot(self.a2, self.W3) + self.b3
        self.a3 = self.sigmoid(self.z3)

        return self.a3

    def compute_loss(self, X, y):
        y_hat = self.predict_proba(X)
        loss = np.mean(0.5 * (y_hat - y)**2)
        return loss

    def backpropagate(self, x, y):
        y_hat = self.forward_propagation(x)

        # Output layer error (MSE derivative * sigmoid derivative)
        # delta3 = (a3 - y) * f'(z3)
        delta3 = (self.a3 - y) * self.sigmoid_derivative(self.a3)

        # Hidden layer 2 error
        # delta2 = (delta3 * W3^T) * f'(z2)
        delta2 = np.dot(delta3, self.W3.T) * self.sigmoid_derivative(self.a2)

        # Hidden layer 1 error
        # delta1 = (delta2 * W2^T) * f'(z1)
        delta1 = np.dot(delta2, self.W2.T) * self.sigmoid_derivative(self.a1)

        # W = W - lr * (input_to_layer^T * delta)
        self.W3 -= self.lr * np.dot(self.a2.T, delta3)
        self.b3 -= self.lr * np.sum(delta3, axis=0, keepdims=True)

        self.W2 -= self.lr * np.dot(self.a1.T, delta2)
        self.b2 -= self.lr * np.sum(delta2, axis=0, keepdims=True)

        self.W1 -= self.lr * np.dot(self.a0.T, delta1)
        self.b1 -= self.lr * np.sum(delta1, axis=0, keepdims=True)

        return np.mean(0.5 * (y_hat - y)**2)

    def fit(self, X_train, y_train, X_val, y_val, max_epochs):
        train_loss_history = []
        val_loss_history = []

        for epoch in range(max_epochs):
            indices = np.arange(X_train.shape[0])
            np.random.shuffle(indices)

            # 1. Train on each sample (SGD)

```



```

        for i in indices:
            self.backpropagate(X_train[i:i+1], y_train[i:i+1])

        # 2. Record costs
        train_loss_history.append(self.compute_loss(X_train, y_train))
        val_loss_history.append(self.compute_loss(X_val, y_val))

        if epoch % 500 == 0:
            print(f"Epoch {epoch}: Train Loss {train_loss_history[-1]:.6f}, Val Loss {val_loss_history[-1]:.6f}")

        return train_loss_history, val_loss_history

def predict_proba(self, X):
    a1 = self.sigmoid(np.dot(X, self.W1) + self.b1)
    a2 = self.sigmoid(np.dot(a1, self.W2) + self.b2)
    y_hat = self.sigmoid(np.dot(a2, self.W3) + self.b3)
    return y_hat

def predict(self, X, decision_thresh=0.5):
    probs = self.predict_proba(X)
    return (probs >= decision_thresh).astype(int)
```

Unit Test

```
In [98]: import unittest
import numpy as np

class TestNeuralNetwork(unittest.TestCase):

    def setUp(self):
        self.n_in = 2
        self.n_h1 = 5
        self.n_h2 = 5
        self.n_out = 1
        self.nn = myNeuralNetwork(self.n_in, self.n_h1, self.n_h2, self.n_out, learning_rate=0.1)

    def test_sigmoid(self):
        test_val = np.array([0])
        expected = 0.5
        self.assertAlmostEqual(self.nn.sigmoid(test_val)[0], expected)

    def test_forward_shape(self):
        x = np.array([[0.5, -0.2]])
        y_hat = self.nn.forward_propagation(x)
        self.assertEqual(y_hat.shape, (1, 1))

    def test_manual_forward(self):
        self.nn.W1 = np.ones((self.n_in, self.n_h1))
        self.nn.b1 = np.zeros((1, self.n_h1))
        self.nn.W2 = np.ones((self.n_h1, self.n_h2))
        self.nn.b2 = np.zeros((1, self.n_h2))
        self.nn.W3 = np.ones((self.n_h2, self.n_out))
        self.nn.b3 = np.zeros((1, self.n_out))

        x = np.array([[1, 1]])
        y_hat = self.nn.forward_propagation(x)

        self.assertGreater(y_hat[0, 0], 0.5)
        self.assertLess(y_hat[0, 0], 1.0)

    def test_backprop_decreases_loss(self):
        x = np.array([[0.1, 0.9]])
        y = np.array([[0.0]])

        initial_prob = self.nn.forward_propagation(x)
        initial_loss = 0.5 * (initial_prob - y)**2

        self.nn.backpropagate(x, y)

        updated_prob = self.nn.forward_propagation(x)
        updated_loss = 0.5 * (updated_prob - y)**2

        self.assertLess(updated_loss, initial_loss)

if __name__ == '__main__':
    unittest.main(argv=['first-arg-is-ignored'], exit=False)
```

```
....
-----
Ran 4 tests in 0.002s

OK
```

Implement XOR learning by my Neural Network

The XOR problem is a classic problem in ML used to demonstrate the necessity of multi-layer neural networks. I put XOR problem in my neural network to look the result.

Initially, the weights (W_1 , W_2 , W_3) were initialized using small random values scaled by a factor of 0.5.

However, the model showed poor performance, with the loss decreasing extremely slowly and the predictions remaining stuck around 0.5.

To overcome this, I increased the weight initialization scale to 1.0.

This adjustment provided the needed for the neurons to operate in a more active range of the Sigmoid function.

As shown in the final Loss Curve, the model was able to break through the initial gradient plateau much faster, achieving a sharp drop in loss and reaching perfect convergence to the target XOR labels [0, 1, 1, 0].

```
In [99]: X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([[0], [1], [1], [0]]) # XOR

nn = myNeuralNetwork(2, 5, 5, 1, learning_rate=0.5)
train_loss, val_loss = nn.fit(X, y, X, y, max_epochs=5000)

print("--- XOR prediction ---")
print(nn.predict(X))
```

```
Epoch 0: Train Loss 0.183693, Val Loss 0.183693
Epoch 500: Train Loss 0.006783, Val Loss 0.006783
Epoch 1000: Train Loss 0.001509, Val Loss 0.001509
Epoch 1500: Train Loss 0.000789, Val Loss 0.000789
Epoch 2000: Train Loss 0.000522, Val Loss 0.000522
Epoch 2500: Train Loss 0.000386, Val Loss 0.000386
Epoch 3000: Train Loss 0.000304, Val Loss 0.000304
Epoch 3500: Train Loss 0.000249, Val Loss 0.000249
Epoch 4000: Train Loss 0.000211, Val Loss 0.000211
Epoch 4500: Train Loss 0.000182, Val Loss 0.000182
--- XOR prediction ---
[[0]
 [1]
 [1]
 [0]]
```

Compare XOR learning result with sklearn neural network

```
In [100]: from sklearn.neural_network import MLPClassifier
import matplotlib.pyplot as plt

# 1. Prepare XOR data
X = np.array([[0,0], [0,1], [1,0], [1,1]])
y = np.array([[0], [1], [1], [0]]) # Reshaped for your model


# 2. Train your own model
nn = myNeuralNetwork(n_in=2, n_layer1=5, n_layer2=5, n_out=1, learning_rate=0.5)
my_train_losses, my_val_losses = nn.fit(X, y, X, y, max_epochs=5000)

# 3. Train sklearn model
# Sklearn expects a 1D array for y
mlp = MLPClassifier(hidden_layer_sizes=(5, 5), activation='logistic', solver='sgd', learning_rate_init=0.5, max_iter=5000, momentum=0, random_state=42, n_iter_no_change=5000, tol=1e-9)
mlp.fit(X, y.ravel())

# 4. Compare Prediction Result
print("---- Compare Prediction Result ----")
print(f"My own model prediction: \n{nn.predict(X).flatten()}")
print(f"Sklearn NN model prediction: \n{mlp.predict(X)}")

# 5. Compare Loss Curve
plt.figure(figsize=(10, 5))
plt.plot(my_train_losses, label='Manual NN (Train) Loss')
plt.plot(mlp.loss_curve_, label='Sklearn MLP Loss', linestyle='--')
plt.title('Loss Curve Comparison (XOR)')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```

Epoch 0: Train Loss 0.135110, Val Loss 0.135110
Epoch 500: Train Loss 0.036754, Val Loss 0.036754
Epoch 1000: Train Loss 0.001807, Val Loss 0.001807
Epoch 1500: Train Loss 0.000761, Val Loss 0.000761
Epoch 2000: Train Loss 0.000463, Val Loss 0.000463
Epoch 2500: Train Loss 0.000328, Val Loss 0.000328
Epoch 3000: Train Loss 0.000251, Val Loss 0.000251
Epoch 3500: Train Loss 0.000202, Val Loss 0.000202
Epoch 4000: Train Loss 0.000169, Val Loss 0.000169
Epoch 4500: Train Loss 0.000144, Val Loss 0.000144
---- Compare Prediction Result ----
My own model prediction:
[0 1 1 0]
Sklearn NN model prediction:
[0 1 1 0]



2.2 Make moons prediction

2.2.1 Data Preparation

I generated 600 samples using make_moons to ensure Ntest = 100 and a total training pool of Ntest = 500.
For hyperparameter tuning, I held out 20% of the training pool as a validation set.
After selecting the optimal parameters, I retrained the final model on the combined 500 samples (Train + Validation) to maximize performance before evaluating on the test set.

```
In [101]: from sklearn.datasets import make_moons
from sklearn.model_selection import train_test_split

# Generate data
X_all, y_all = make_moons(n_samples=600, noise=0.20, random_state=42)
y_all = y_all.reshape(-1, 1)

# Split: train+val
X_train_val_raw, X_test, y_train_val_raw, y_test = train_test_split(X_all, y_all, test_size=100, random_state=42)

# Split validation set from train+val
X_train, X_val, y_train, y_val = train_test_split(X_train_val_raw, y_train_val_raw, test_size=0.2, random_state=42)
```

2.2.2 Tracking both training and validation loss per epoch to create the cost func plots

```
In [102]: def fit(self, X_train, y_train, X_val, y_val, max_epochs):
    train_loss_history = []
    val_loss_history = []

    for epoch in range(max_epochs):
        # 1. Stochastic Gradient Descent
        indices = np.arange(X_train.shape[0])
        np.random.shuffle(indices)
        for i in indices:
            # Update weights using one sample at a time
            self.backpropagate(X_train[i:i+1], y_train[i:i+1])

        # 2. Record costs for the "Cost function plots" requirement
        train_loss_history.append(self.compute_loss(X_train, y_train))
        val_loss_history.append(self.compute_loss(X_val, y_val))

        if epoch % 500 == 0:
            print(f"Epoch {epoch}: Train Loss {train_loss_history[-1]:.4f}, Val Loss {val_loss_history[-1]:.4f}")

    return train_loss_history, val_loss_history
```

2.2.3 Hyperparameter Tuning

```
In [103]: best_val_loss = float('inf')
best_params = {}

for lr in [0.1, 0.5, 1.0]:
    for epochs in [1000, 2000, 5000]:
        nn = myNeuralNetwork(2, 5, 5, 1, learning_rate=lr)
        # Call with positional arguments matches the new signature
        train_h, val_h = nn.fit(X_train, y_train, X_val, y_val, max_epochs=epochs)

        if val_h[-1] < best_val_loss:
            best_val_loss = val_h[-1]
            best_params = {'lr': lr, 'epochs': epochs}

print(f"Best Hyperparameters found: {best_params}")
```

```
Epoch 0: Train Loss 0.125513, Val Loss 0.124777
Epoch 500: Train Loss 0.043498, Val Loss 0.050117
Epoch 0: Train Loss 0.103355, Val Loss 0.096518
Epoch 500: Train Loss 0.012419, Val Loss 0.019519
Epoch 1000: Train Loss 0.012030, Val Loss 0.022167
Epoch 1500: Train Loss 0.011849, Val Loss 0.023498
Epoch 0: Train Loss 0.113650, Val Loss 0.110957
Epoch 500: Train Loss 0.012781, Val Loss 0.021106
Epoch 1000: Train Loss 0.012693, Val Loss 0.019933
Epoch 1500: Train Loss 0.011724, Val Loss 0.022916
Epoch 2000: Train Loss 0.011797, Val Loss 0.021503
Epoch 2500: Train Loss 0.011637, Val Loss 0.025443
Epoch 3000: Train Loss 0.010717, Val Loss 0.024306
Epoch 3500: Train Loss 0.010436, Val Loss 0.024743
Epoch 4000: Train Loss 0.010237, Val Loss 0.024732
Epoch 4500: Train Loss 0.010103, Val Loss 0.024647
Epoch 0: Train Loss 0.063034, Val Loss 0.056172
Epoch 500: Train Loss 0.012287, Val Loss 0.023586
Epoch 0: Train Loss 0.081222, Val Loss 0.075199
Epoch 500: Train Loss 0.014505, Val Loss 0.019109
Epoch 1000: Train Loss 0.012107, Val Loss 0.023787
Epoch 1500: Train Loss 0.011335, Val Loss 0.023365
Epoch 0: Train Loss 0.067199, Val Loss 0.059783
Epoch 500: Train Loss 0.011837, Val Loss 0.022267
Epoch 1000: Train Loss 0.011205, Val Loss 0.024973
Epoch 1500: Train Loss 0.010836, Val Loss 0.025486
Epoch 2000: Train Loss 0.010319, Val Loss 0.026074
Epoch 2500: Train Loss 0.009790, Val Loss 0.025592
Epoch 3000: Train Loss 0.010327, Val Loss 0.025697
Epoch 3500: Train Loss 0.009067, Val Loss 0.027237
Epoch 4000: Train Loss 0.009332, Val Loss 0.026310
Epoch 4500: Train Loss 0.009017, Val Loss 0.027105
Epoch 0: Train Loss 0.047738, Val Loss 0.045649
Epoch 500: Train Loss 0.011320, Val Loss 0.025362
Epoch 0: Train Loss 0.052797, Val Loss 0.051583
Epoch 500: Train Loss 0.011615, Val Loss 0.023735
Epoch 1000: Train Loss 0.010690, Val Loss 0.024709
Epoch 1500: Train Loss 0.010139, Val Loss 0.025609
Epoch 0: Train Loss 0.050066, Val Loss 0.050628
Epoch 500: Train Loss 0.012263, Val Loss 0.023470
Epoch 1000: Train Loss 0.011179, Val Loss 0.024503
Epoch 1500: Train Loss 0.009266, Val Loss 0.027364
Epoch 2000: Train Loss 0.008970, Val Loss 0.032483
Epoch 2500: Train Loss 0.008477, Val Loss 0.039783
Epoch 3000: Train Loss 0.008142, Val Loss 0.036486
Epoch 3500: Train Loss 0.008161, Val Loss 0.041172
Epoch 4000: Train Loss 0.007569, Val Loss 0.035936
Epoch 4500: Train Loss 0.007652, Val Loss 0.038076
Best Hyperparameters found: {'lr': 0.5, 'epochs': 1000}
```

The best hyperparameter is learning rate: 0.1 with 1000 epochs

Regarding the selection of training epochs, I observed that 1,000 epochs provided the optimal balance between learning and generalization.

While the training loss continued to decrease slightly as the epochs increased to 5,000, the validation loss began to plateau or even rise marginally.

2.2.4 Cost function plots


This plot helps visualize the learning progress and detect overfitting.

In [104]: `import matplotlib.pyplot as plt`

```
# Retrain using the best hyperparameters found: lr=0.1, epochs=1000
nn_best = myNeuralNetwork(n_in=2, n_layer1=5, n_layer2=5, n_out=1, learning_rate=0.1)
train_h, val_h = nn_best.fit(X_train, y_train, X_val, y_val, max_epochs=1000)

plt.figure(figsize=(8, 5))
plt.plot(train_h, label='Training Cost')
plt.plot(val_h, label='Validation Cost')
plt.title('Cost Function Curves (lr=0.1, epochs=1000)')
plt.xlabel('Epochs')
plt.ylabel('MSE Loss')
plt.legend()
plt.grid(True)
plt.show()
```

Epoch 0: Train Loss 0.132470, Val Loss 0.129569
Epoch 500: Train Loss 0.013110, Val Loss 0.018709



2.2.5 Decision boundary plots


In [105]: `import numpy as np`

```
def plot_boundary(nn, X, y, ax, title):
    # Set min and max values and give it some padding
    x_min, x_max = X[:, 0].min() - 0.5, X[:, 0].max() + 0.5
    y_min, y_max = X[:, 1].min() - 0.5, X[:, 1].max() + 0.5
    h = 0.02
    # Generate a grid of points with distance h between them
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    # Predict the function value for the whole grid
    Z = nn.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot the contour and training examples
    ax.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.Spectral)
    ax.scatter(X[:, 0], X[:, 1], c=y.ravel(), s=40, cmap=plt.cm.Spectral, edgecolors='k')
    ax.set_title(title)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 5))
plot_boundary(nn_best, X_train, y_train, ax1, "Training Data Decision Boundary")
plot_boundary(nn_best, X_val, y_val, ax2, "Validation Data Decision Boundary")
plt.show()
```



The visualization confirms that your two-layer architecture (2-5-5-1) has sufficient capacity to learn the non-linear moons.

- Training Subplot: The boundary cleanly curves out the two crescent shapes, handling the overlap areas logically.
- Validation Subplot: The boundary holds up well against the validation data, which proves the model didn't just memorize the training points but actually learned the underlying distribution.

ROC Curve and AUC Score

```
In [106]: from sklearn.metrics import roc_curve, auc
from sklearn.neural_network import MLPClassifier

# Retrain on combined dataset (Train + Validation)
X_train_val = np.vstack((X_train, X_val))
y_train_val = np.vstack((y_train, y_val))


# Final Manual Model training
nn_final = myNeuralNetwork(2, 5, 5, 1, learning_rate=0.1)
_ = nn_final.fit(X_train_val, y_train_val, X_train_val, y_train_val, max_epochs=1000)
y_prob_manual = nn_final.predict_proba(X_test)

# Sklearn Comparison Model
mlp = MLPClassifier(hidden_layer_sizes=(5, 5), activation='logistic', solver='sgd',
                    learning_rate_init=0.1, max_iter=1000, random_state=42)
mlp.fit(X_train_val, y_train_val.ravel())
y_prob_sklearn = mlp.predict_proba(X_test)[:, 1]

# Calculate ROC and AUC
fpr_m, tpr_m, _ = roc_curve(y_test, y_prob_manual)
fpr_s, tpr_s, _ = roc_curve(y_test, y_prob_sklearn)

plt.figure(figsize=(8, 6))
plt.plot(fpr_m, tpr_m, label=f'Manual NN (AUC = {auc(fpr_m, tpr_m):.2f})')
plt.plot(fpr_s, tpr_s, label=f'Sklearn MLP (AUC = {auc(fpr_s, tpr_s):.2f})')
plt.plot([0, 1], [0, 1], 'k--', label='Chance Diagonal') # Include chance diagonal
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison (Test Data)')
plt.legend()
plt.show()
```

Epoch 0: Train Loss 0.085375, Val Loss 0.085375
Epoch 500: Train Loss 0.014166, Val Loss 0.014166



ROC Curve Comparison

- Manual NN (AUC = 0.99): This implementation reached near-perfect classification performance on the test set.
- Sklearn MLP (AUC = 0.95): Manual model actually outperformed the scikit-learn baseline on this specific test run.

2.4 Two-ways to improve my model

1. ReLU: Replacing the sigmoid function in the hidden layer with ReLU would improve the network in sparsity and efficiency. Because ReLU effectively turn off neurons with negative inputs, leading to a sparser and more computationally efficient representation.
2. Improved weight initialization: Use heuristic based initialization like Xavier. Because setting weights to 0 is bad because it leads too much symmetry, causing many gradients to be identical. Xavier ensures that weights are neither too small nor too big.